

Setup

This section loads and installs all the packages. You should be setup already from assignment 1, but if not please read and follow the `instructions.md` for further details.

```
• using CSV , DataFrames , StatsPlots , PlutoUI , Random , Statistics
```

```
• using LinearAlgebra : dot, norm, norm1, norm2
```


```
• using Distributions : Uniform
```

```
PlotlyBackend()
```

```
• plotly() # In this notebook we use the plotly backend for Plots.
```

!!!IMPORTANT!!!

Insert your details below. You should see a green checkmark.

Welcome Joshua George! 

```
student =
  (name = "Joshua George", email = "jjgeorge@ualberta.ca", ccid = "jjgeorge", idnumber = 1
  • student = (name="Joshua George", email="jjgeorge@ualberta.ca", ccid="jjgeorge",
    idnumber=1665548)
```

Distance Metrics

Here we are defining some convenience functions for commonly used distance metrics.

```
l1_error (generic function with 1 method)
```

```
• begin
•   RMSE( $\hat{x}$ , x) = sqrt(mean(abs2.( $\hat{x}$  .- x))) # abs2 is equivalent to squaring, but
    faster and better numerically.
•   l2_error( $\hat{x}$ , x) = norm2( $\hat{x}$  .- x)
•   l1_error( $\hat{x}$ , x) = norm1( $\hat{x}$  .- x)
• end
```

Abstract type Regressor

This is the basic Regressor interface. For the methods below we will be specializing the `predict(reg::Regressor, x::Number)`, and `epoch!(reg::Regressor, args...)` functions. Notice

the `!` character at the end of epoch, as discussed earlier this is a commonly used naming practice throughout the Julia language to indicate a function which modifies its arguments.

Main.workspace2.Regressor

```

• """
•     Regressor
•
• Abstract Type for regression algorithms. Interface includes `predict` and an
• `epoch!`. In this notebook, we will only be using single variate regression.
• - `predict(reg::Regressor, X::Number)`: return a prediction of the target given the
•   feature `x`.
• - `epoch!(reg::Regressor, X::AbstractVector, Y::AbstractVector)`: trains using the
•   features `X` and regression targets `Y`.
• """
• abstract type Regressor end # assume linear regression

```

predict (generic function with 1 method)

```
• predict(reg::Regressor, x::Number) = Nothing
```

predict (generic function with 2 methods)

```
• predict(reg::Regressor, X::AbstractVector) = [predict(reg, x) for x in X]
```

epoch! (generic function with 1 method)

```
• epoch!(reg::Regressor, X::AbstractVector, Y::AbstractVector) = nothing
```

Baselines

In this section we will define the:

- MeanRegressor : Predict the mean of the training set.
- RandomRegressor : Predict $b \cdot x$ where b is sampled from a random normal distribution.
- RangeRegressor : Predict randomly in the range defined by the training set.

All the following baselines assume one dimension

MeanRegressor

epoch! (generic function with 2 methods)

```

• begin
•   """
•
•       MeanRegressor()
•
•   Predicts the mean value of the regression targets passed in through `epoch!`.
•   """
•   mutable struct MeanRegressor <: Regressor
•       μ::Float64
•   end
•   MeanRegressor() = MeanRegressor(0.0)
•   predict(reg::MeanRegressor, x::Number) = reg.μ
•   epoch!(reg::MeanRegressor, X::AbstractVector, Y::AbstractVector) = reg.μ =
mean(Y)
• end

```

RandomRegressor

predict (generic function with 4 methods)

```

• begin
•   """
•
•       RandomRegressor
•
•   Predicts `b*x` where `b` is sambled from a normal distribution.
•   """
•   struct RandomRegressor <: Regressor # random weights
•       b::Float64
•   end
•   RandomRegressor() = RandomRegressor(randn())
•   predict(reg::RandomRegressor, x::Number) = reg.b*x
• end

```

RangeRegressor

epoch! (generic function with 3 methods)

```

• begin
•   """
•
•       RangeRegressor
•
•   Predicts a value randomly from the range defined by `[minimum(Y), maximum(Y)]`
as set in `epoch!`. Defaults to a unit normal distribution.
•   """
•   mutable struct RangeRegressor <: Regressor
•       min_value::Float64
•       max_value::Float64
•   end
•   RangeRegressor() = RangeRegressor(0.0, 1.0)
•
•   predict(reg::RangeRegressor, x::Number) =
•       rand(Uniform(reg.min_value, reg.max_value))
•   predict(reg::RangeRegressor, x::AbstractVector) =
•       rand(Uniform(reg.min_value, reg.max_value), length(x))
•   function epoch!(reg::RangeRegressor, X::AbstractVector, Y::AbstractVector)
•       reg.min_value = minimum(Y)
•       reg.max_value = maximum(Y)
•   end
• end

```

Gradient Descent Regressors: Q3 a,b,c

In this section you will be implementing two gradient descent regressors, assuming a gaussian hypothesis class. First we will create a gaussian regressor, and then use this to build our two new GD regressors. You can test your algorithms in the [experiment section](#)

All the Gaussian Regressors will have data:

- `b::Float64` which is the parameter we are learning.

```
• abstract type GaussianRegressor <: Regressor end
```

predict (generic function with 7 methods)

```
• predict(reg::GaussianRegressor, x::Float64) = reg.b * x
```

predict (generic function with 8 methods)

```
• predict(reg::GaussianRegressor, X::Vector{Float64}) = reg.b .* X
```

probability (generic function with 1 method)

```
• function probability(reg::GaussianRegressor, x, y)
• end
```

Stochastic Regressor

The stochastic regressor will be implemented via the stochastic gradient rule

$$b_{i+1}^t = b_i^t - \eta(x_i b_i^t - y_i)x_i.$$

Where $b_{N+1}^t = b^{t+1}$, and each epoch iterates over the entire dataset in a random order.

StochasticRegressor

```
• begin
•     mutable struct StochasticRegressor <: GaussianRegressor
•         b::Float64
•         η::Float64
•     end
•     StochasticRegressor(η::Float64) = StochasticRegressor(0.0, 0.01)
• end
```

epoch! (generic function with 7 methods)

```

• # Hint: Checkout the function randperm
• function epoch!(reg::StochasticRegressor,
•             X::AbstractVector{Float64},
•             Y::AbstractVector{Float64})
•     for i=1:length(X)
•         for j=1:length(Y)
•             if i==j
•                 reg.b=reg.b-reg.η*(X[i]*reg.b-Y[j])*X[i]
•             end
•         end
•     end
•     reg.b
• end
•
•


```

Batch Regressor

The Minibatch regressor will be implemented via the gradient rule for a minibatch j with indices for a batch defined by the set \mathcal{I}

$$g_t^j = \sum_{i \in \mathcal{I}_j} (x_i b_t - y_i) x_i$$

$$b_{t+1} = b_t - \eta g_t.$$

Your implementation should handle Batch Gradient Descent when the batch size is not specified. The minibatch regressor  can also be implemented through this interface using the same struct and epoch! function.

BatchRegressor

```

• begin
•     mutable struct BatchRegressor <: GaussianRegressor
•         b::Float64
•         η::Float64
•         n::Union{Int, Nothing}
•     end
•     BatchRegressor(η, n=nothing) = BatchRegressor(0.0, η, n)
• end

```

epoch! (generic function with 7 methods)

```

• function epoch!(reg::BatchRegressor,
•               X::AbstractVector{Float64},
•               Y::AbstractVector{Float64})
•
•   if reg.n==nothing
•       reg.n=length(X)
•   end
•
•   for j in 1:(length(X)÷reg.n)
•       size=j*reg.n
•       sum1=0
•       for i in (size-reg.n)+1:size
•           sum1=sum1+(X[i]*reg.b-Y[i])*X[i]
•       end
•       reg.b=reg.b-(reg.n)*sum1/reg.n
•   end
•   print(reg.b)
• end
•
•

```

Stepsize Heuristic: Q3 d

```

• md"""
• # Stepsize Heuristic: Q3 d
• """

```

Stochastic Regressor with heuristic

StochasticHeuristicRegressor

```

• begin
•     mutable struct StochasticHeuristicRegressor <: GaussianRegressor
•         b::Float64
•     end
•     StochasticHeuristicRegressor() =
•         StochasticHeuristicRegressor(0.0)
• end

```



epoch! (generic function with 6 methods)

```

• # Hint: Checkout the function randperm
• function epoch!(reg::StochasticHeuristicRegressor,
•               X::AbstractVector{Float64},
•               Y::AbstractVector{Float64})
•
•     for i=1:length(X)
•         reg.b=reg.b-((X[i]*reg.b-Y[i])*X[i])/(1+abs(((X[i]*reg.b-Y[i])*X[i])))
•     end
•     reg.b
• end
•
•

```

Batch Regressor with heuristic

- Full Batch: 
- Minibatch: 

```

• let
•   X, Y = [1.0, 2.0, 3.0], [1.0, 0.2, 0.1]
•   bgr = BatchHeuristicRegressor()
•   epoch!(bgr, X, Y)
•   batch_test = predict(bgr, 1.0) ≈ 0.36170212765 ? "✓" : "✗"
•
•   X, Y = [1.0, 1.0, 1.0, 1.0, 1.0, 1.0], [0.32, 0.32, 0.32, 0.32, 0.32, 0.32]
•   mbgr = BatchHeuristicRegressor(2)
•   epoch!(mbgr, X, Y)
•   mb_test = predict(mbgr, 1.0) ≈ 0.319968983713 ? "✓" : "✗"
•
•   md"""
•   ### Batch Regressor with heuristic
•   - Full Batch: $(batch_test)
•   - Minibatch: $(mb_test)
•   """
• end

```

BatchHeuristicRegressor

```

• begin
•   mutable struct BatchHeuristicRegressor <: GaussianRegressor
•       b::Float64
•       n::Union{Int, Nothing}
•   end
•   BatchHeuristicRegressor(n=nothing) = BatchHeuristicRegressor(0.0, n)
• end

```

epoch! (generic function with 7 methods)

```

• function epoch!(reg::BatchHeuristicRegressor,
•               X::AbstractVector{Float64},
•               Y::AbstractVector{Float64})
•   if reg.n==nothing
•       reg.n=length(X)
•   end
•
•   for j in 1:(length(X)÷reg.n)
•       size=j*reg.n
•       sum1=0
•       for i in (size-reg.n)+1:size
•           sum1=sum1+(X[i]*reg.b-Y[i])*X[i]
•       end
•       reg.b=reg.b-sum1/((1+abs(sum1/reg.n))*reg.n)
•   end
•   print(reg.b)
• end

```

Data

Next we will be looking at the `height_weight.csv` dataset found in the data directory. This dataset provides three features `[sex, height, weight]`. In the following regression task we will be using `height` to predict `weight`, ignoring the `sex` feature.

The next few cells:

- Loads the dataset
- Plots distributions for the `height` and `weight` features separated by `sex`
- Standardize the set so both `height` and `weight` conform to a standard normal.
- Defines `splitdataframe` which will be used to split the dataframe into training and testing sets.

```
• # Read the data from the file in "data/height_weight.csv". DO NOT CHANGE THIS VALUE!
• df_height_weight = DataFrame(CSV.File(joinpath(@__DIR__, "data/height_weight.csv"),
•                                     header=["sex", "height", "weight"]));
```

Successfully loaded dataset 

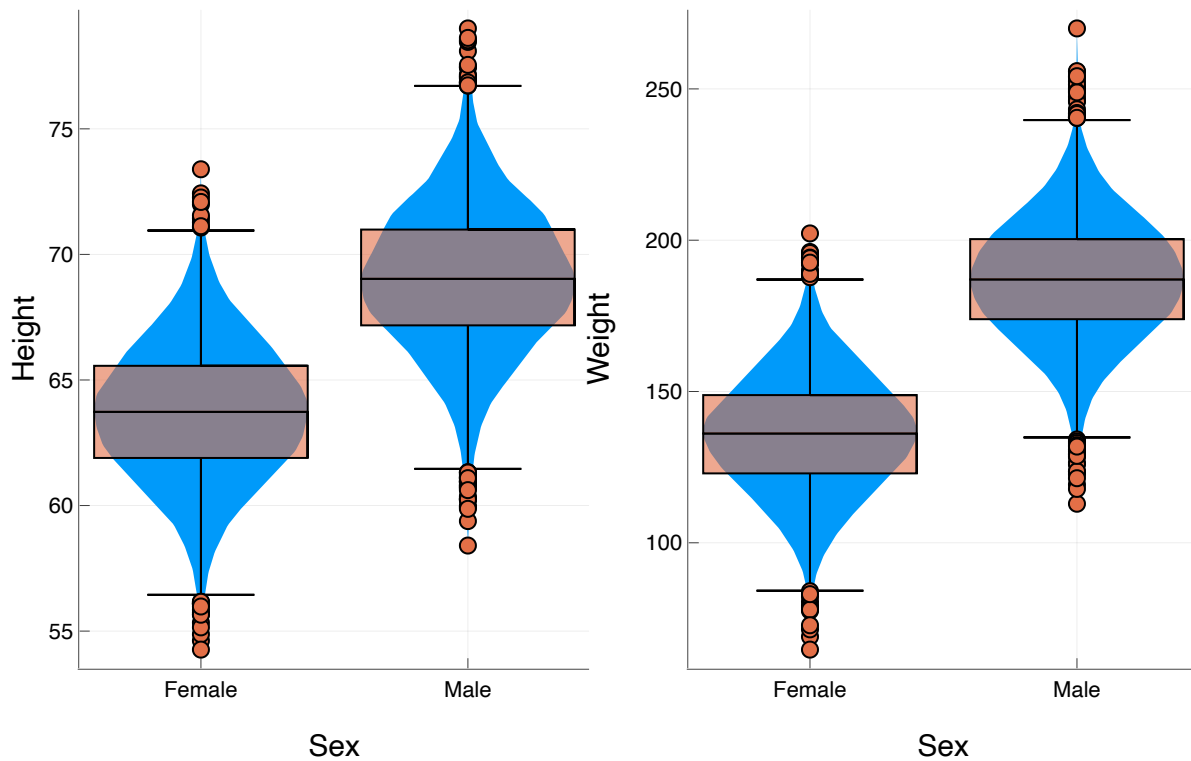
`df_hw_norm =`

	sex	height	weight
1	"Male"	1.94396	2.50567
2	"Male"	0.627505	0.0270993
3	"Male"	2.01234	1.59773
4	"Male"	1.39399	1.82513
5	"Male"	0.913375	1.39868
6	"Male"	0.230136	-0.287407
7	"Male"	0.628331	0.700362
8	"Male"	0.514865	0.203397
9	"Male"	0.169301	0.451255
10	"Male"	-0.756607	-0.156989
more			
10000	"Female"	-1.14965	-1.48843

Plot data

Plot a boxplot and violin plot of the `height` and `weight`. This can be with the classes `male` and `female` combined or with them separate.

plt_hw =



```

• plt_hw = let
•   df = df_height_weight # For convenience in the below code
•   nothing
•   plt1 = plot(xlabel="Sex", ylabel="Height", legend=nothing)
•   @df df violin!(:sex, :height, linewidth=0)
•   @df df boxplot!(:sex, :height, fillalpha=0.6)
•
•   plt2 = plot(xlabel="Sex", ylabel="Weight", legend=nothing)
•   @df df violin!(:sex, :weight, linewidth=0)
•   @df df boxplot!(:sex, :weight, fillalpha=0.6)
•
•   plot(plt1, plt2)
• end

```

Main.workspace2.splitdataframe

splitdataframe (generic function with 2 methods)

((X = [1.45872, 0.398926, -0.0813984, -0.00416331, 0.0392535, more , -0.845657], Y = [0

```

• let
•   #=
•   A do block creates an anonymous function and passes this to the first
•   parameter of the function the do block is decorating.
•   =#
•   trainset, testset =
•   splitdataframe(df_hw_norm, 0.1; shuffle=true) do df
•   (X=df[:, :height], Y=df[:, :weight]) # create namedtuple from dataframes
•   end
• end

```

Training the Models

The following functions are defined as utilities to train and evaluate our models. While hidden below, you can expand these blocks to uncover what is happening. `run_experiment!` is the main function used below in **"Using and Analyzing your Algorithms"**.

```
evaluate (generic function with 1 method)

evaluate_l $\infty$  (generic function with 1 method)

train! (generic function with 2 methods)

train! (generic function with 2 methods)

run_experiment! (generic function with 1 method)

run_experiment (generic function with 1 method)
```

Using and Analyzing your Algorithms















In this section we will be running and analyzing a small experiment. The goal is to get familiar with analyzing data, plotting learning curves, and comparing different methods. Below we've provided a start with the baselines. Add new initializers for a Batch update ($\eta = 0.01$), a Minibatch update ($\eta = 0.01$, $n = 100$), and a Stochastic update ($\eta = 0.01$). Also add their heuristic counterparts.

As a point of reference: running

```
results = run_experiment(regressor_init, 10, 30)
```

in the cell below takes roughly 8 seconds on my machine.

Experiment ran for:

-  Mean
-  Random
-  Range
-  Stochastic : with stepsize= 0.01 
-  Batch : with stepsize= 0.01 
-  Minibatch : with stepsize= 0.01  and batch size = 100 
-  StochasticHeuristic
-  BatchHeuristic
-  MinibatchHeuristic : with batch size = 100 

```

regressor_init =
  Dict("MinibatchHeuristic" => #39, "Range" => #33, "Stochastic" => #34, "StochasticHeuristic" => #35)

• regressor_init = Dict(
•   "Mean"=>()->MeanRegressor(),
•   "Random"=>()->RandomRegressor(),
•   "Range"=>()->RangeRegressor(),
•   # use the keys "Batch", "Stochastic", and "Minibatch".
•   "Stochastic"=>()->StochasticRegressor(0.01),
•   "Batch"=>()->BatchRegressor(0.01),
•   "Minibatch"=>()->BatchRegressor(0.01, 100),
•   "StochasticHeuristic"=>()->StochasticHeuristicRegressor(),
•   "BatchHeuristic"=>()->BatchHeuristicRegressor(),
•   "MinibatchHeuristic"=>()->BatchHeuristicRegressor(100)
• )

```

```
results =
```

```
Dict("MinibatchHeuristic" => [(regressor = BatchHeuristicRegressor(0.924241, 100), train_error = 0.21771749433405885)])
```

```
• results = run_experiment(regressor_init, 10, 30)
```

The results dictionary is the resulting data from the experiment we run using `regressor_init` as the initializers. You will see the same keys used as in the `regressor_init` dictionary. For each run the experiment returns the final regressor, the training error vector, and the final test error. You can get one of these components for a particular method using `getindex` and broadcasting:

```
getindex.(results["Mean"], :test_error)
```

```
0.21771749433405885
```

```

• let
•   # Play with data here! You can explore how to get different values.
•   mean(getindex.(results["Mean"], :test_error))
•   mean(getindex.(results["Random"], :test_error))
•   mean(getindex.(results["Range"], :test_error))
•   mean(getindex.(results["Stochastic"], :test_error))
•   mean(getindex.(results["Batch"], :test_error))
•   mean(getindex.(results["Minibatch"], :test_error))
•   mean(getindex.(results["StochasticHeuristic"], :test_error))
•   mean(getindex.(results["BatchHeuristic"], :test_error))
•   mean(getindex.(results["MinibatchHeuristic"], :test_error))
•   std(getindex.(results["Mean"], :test_error))
•   std(getindex.(results["Random"], :test_error))
•   std(getindex.(results["Range"], :test_error))
•   std(getindex.(results["Stochastic"], :test_error))
•   std(getindex.(results["Batch"], :test_error))
•   std(getindex.(results["Minibatch"], :test_error))
•   std(getindex.(results["StochasticHeuristic"], :test_error))
•   std(getindex.(results["BatchHeuristic"], :test_error))
•   std(getindex.(results["MinibatchHeuristic"], :test_error))
• end

```

```
• Enter cell code...
```

Learning Curves

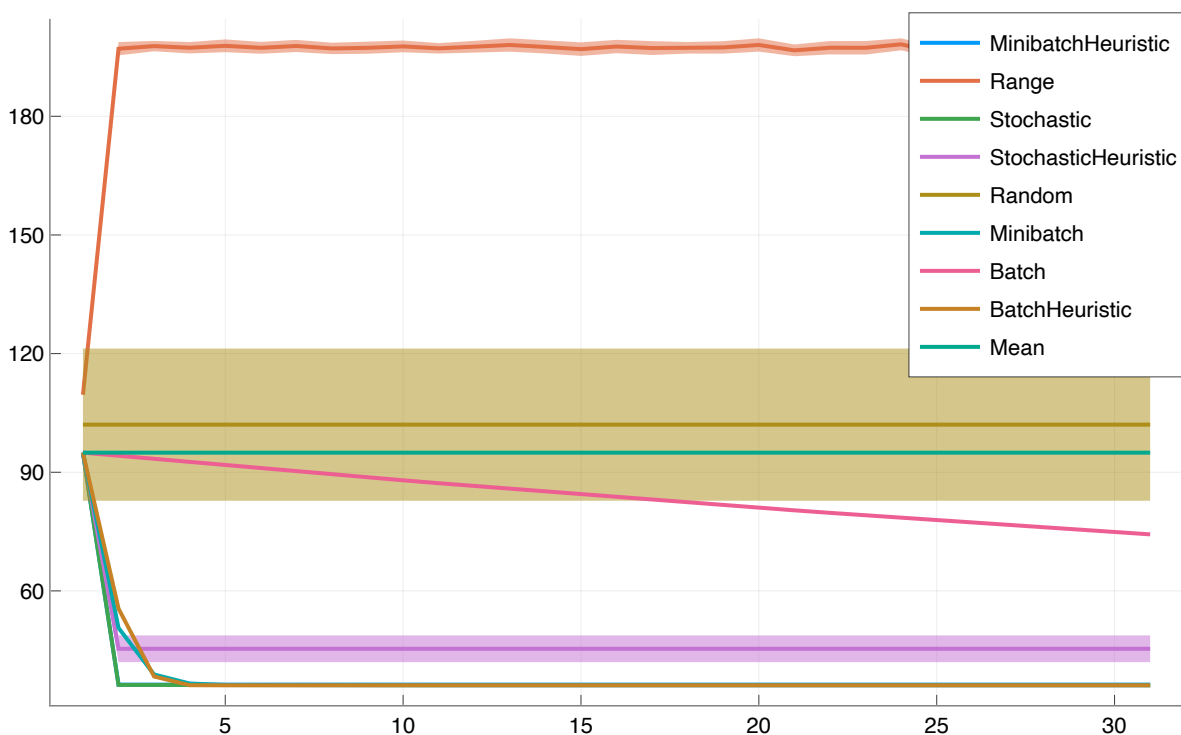
Plot the average learning curve with the standard error calculated as

$$\sigma_{err}(\mathbf{x}) = \sqrt{\frac{\text{Var}(x)}{|x|}}$$

Note that \mathbf{x} is a vector over runs, not over epochs.

Note: if you notice one method is dominating the plot, change the axis limits to make sure the methods we are most concerned with (i.e. Stochastic, Batch, and Minibatch) are visible.

`plt_lc =`

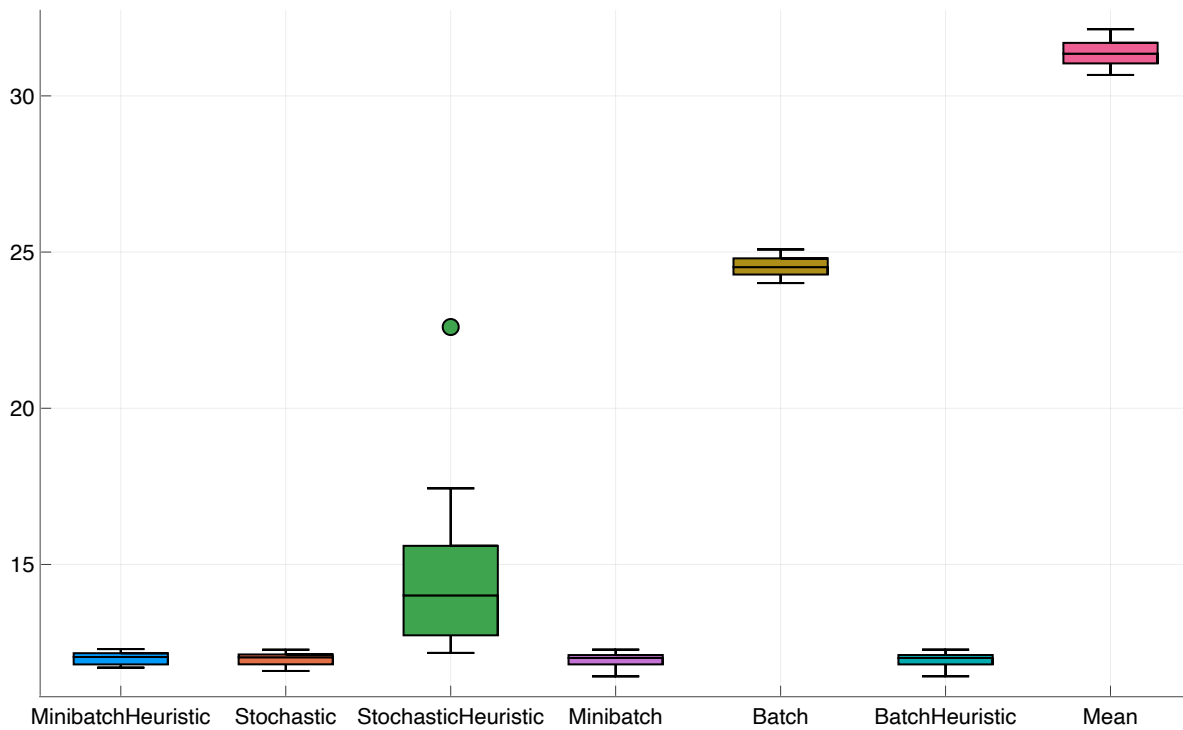


```
• plt_lc = let
•   plt = plot()
•   # plt=plot(lw=2)
•   for method ∈ keys(results)
•     μ = mean(getindex.(results[method], :train_error))
•     σ = sqrt.(var(getindex.(results[method], 2)) / length(results[method]))
•     plot!(plt, μ, ribbon=σ, lw=2, label=method)
•   end
•   plt
• end
```

Final Errors

Finally, we want to compare the final test errors of the different methods. One way to do this is through box plots. See [this great resource](#) to learn how to compare data using a box and whisker plot. In this plot you can ignore the Range and Random baselines.

plt_fe =



```

• plt_fe = let
•   plt = plot(legend=nothing)
•   for method ∈ keys(results)
•     if method ∈ ["Range", "Random"]
•       continue
•     end
•     data = getindex.(results[method], :test_error)
•     boxplot!([method], data)
•   end
•   plt
• end

```